

## HW 10.1

```
rm(list=ls())

crimedata <- read.table("C:\\Users\\Adnan Karim\\Documents\\ISYE 6501\\Homeworks\\hw6-SP22\\data
9.1\\uscrime.txt", stringsAsFactors = F, header = T)

head(crimedata)
```

```
##      M So   Ed Po1  Po2   LF   M.F Pop   NW   U1  U2 Wealth Ineq   Prob
## 1 15.1  1  9.1  5.8  5.6 0.510 95.0 33 30.1 0.108 4.1  3940 26.1 0.084602
## 2 14.3  0 11.3 10.3  9.5 0.583 101.2 13 10.2 0.096 3.6  5570 19.4 0.029599
## 3 14.2  1  8.9  4.5  4.4 0.533 96.9 18 21.9 0.094 3.3  3180 25.0 0.083401
## 4 13.6  0 12.1 14.9 14.1 0.577 99.4 157 8.0 0.102 3.9  6730 16.7 0.015801
## 5 14.1  0 12.1 10.9 10.1 0.591 98.5 18 3.0 0.091 2.0  5780 17.4 0.041399
## 6 12.1  0 11.0 11.8 11.5 0.547 96.4 25 4.4 0.084 2.9  6890 12.6 0.034201
##      Time Crime
## 1 26.2011    791
## 2 25.2999   1635
## 3 24.3006    578
## 4 29.9012   1969
## 5 21.2998   1234
## 6 20.9995    682
```

```
library(tree)

#here i am importing the data into R and displaying it
#now i will build the regression tree model

crimedataatree <- tree(Crime~., data = crimedata)

summary(crimedataatree)
```


```
##
## Regression tree:
## tree(formula = Crime ~ ., data = crimedata)
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "LF" "NW"
## Number of terminal nodes: 7
## Residual mean deviance: 47390 = 1896000 / 40
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.900 -98.300  -1.545   0.000 110.600  490.100
```

```
#we can now use the frame callout to see how this tree was split up

crimedataatree$frame
```

##	var	n	dev	yval	splits.cutleft	splits.cutright
## 1	Po1	47	6880927.66	905.0851	<7.65	>7.65
## 2	Pop	23	779243.48	669.6087	<22.5	>22.5
## 4	LF	12	243811.00	550.5000	<0.5675	>0.5675
## 8	<leaf>	7	48518.86	466.8571		
## 9	<leaf>	5	77757.20	667.6000		
## 5	<leaf>	11	179470.73	799.5455		
## 3	NW	24	3604162.50	1130.7500	<7.65	>7.65
## 6	Pop	10	557574.90	886.9000	<21.5	>21.5
## 12	<leaf>	5	146390.80	1049.2000		
## 13	<leaf>	5	147771.20	724.6000		
## 7	Po1	14	2027224.93	1304.9286	<9.65	>9.65
## 14	<leaf>	6	170828.00	1041.0000		
## 15	<leaf>	8	1124984.88	1502.8750		

```
plot(crimedatatree)
text(crimedatatree)
```




```
#we can now prune this tree by setting the nodes to 5

nodes <- 5

crimedatatreeprune <- prune.tree(crimedatatree, best = nodes)

plot(crimedatatreeprune)
text(crimedatatreeprune)
```



```
summary(crimedatatreeprune)
```

```
##
## Regression tree:
## snip.tree(tree = crimedatatree, nodes = c(4L, 6L))
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "NW"
## Number of terminal nodes: 5
## Residual mean deviance: 54210 = 2277000 / 42
## Distribution of residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.9 -107.5   15.5     0.0  122.8   490.1
```

```
#interpreting this data shows that 5 nodes is leading to more error and overfitting in this model, the original 7 nodes works just fine
```

```
nodes2 <- 7
```

```
crimedatatreeprune2 <- prune.tree(crimedatatree, best = nodes2)
```

```
plot(crimedatatreeprune2)
text(crimedatatreeprune2)
summary(crimedatatreeprune2)
```

```
##
## Regression tree:
## tree(formula = Crime ~ ., data = crimedata)
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "LF" "NW"
## Number of terminal nodes: 7
## Residual mean deviance: 47390 = 1896000 / 40
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.900 -98.300  -1.545   0.000 110.600  490.100
```

```
#now we can calculate r2 to see the quality of this model
```

```
crimedataprediction <- predict(crimedatatreeprune2, data= crimedata[,1:15])
rss <- sum((crimedataprediction - crimedata[,16])^2)
tss <- sum((crimedata[,16] - mean(crimedata[,16]))^2)

r2 <- 1- rss/tss

r2
```

```
## [1] 0.7244962
```

```
# the r2 at 0.724 is not the best, so even this model has a degree of over fitting
```

```
# now we can do a random forest model and compare
```

```
library(randomForest)
```

```
## randomForest 4.7-1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```

crimforest <- randomForest(Crime~. , data = crimedata, importance = T, nodesize = 5)

crimeforpredict <- predict(crimforest, data = crimedata[,-16])

rss2 <- sum((crimeforpredict - crimedata[,16])^2)

r2new <- 1 - rss2/tss

r2new

```

```
## [1] 0.4417235
```

*# the r2 value with the forest method is even worse than the single tree , so it is not as accurate*

```
#### 10.2
```

*#I can use a logistic regression model in my job to predict the probability that a certain price of a product (chemical) would sell in the open market. So things like volume, quality, market price could all be used as predictors.*

```
#### 10.3
```

```
gcred <- read.table("C:\\Users\\Adnan Karim\\Documents\\ISYE 6501\\Homeworks\\hw7-SP22\\data 10.3\\germancredit.txt", stringsAsFactors = F, header = F)
```

```
library(caret)
```

```
## Loading required package: ggplot2
```


```
## Warning in register(): Can't find generic `scale_type` in package ggplot2 to
## register S3 method.
```

```
##
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:randomForest':
```

```
##
## margin
```

```
## Loading required package: lattice
```



*#caret is so we can use the data partition function later*

`head(gcred)`

```
##      V1 V2  V3  V4   V5  V6  V7 V8  V9  V10 V11  V12 V13  V14  V15 V16  V17 V18
## 1 A11  6 A34 A43 1169 A65 A75  4 A93 A101  4 A121  67 A143 A152  2 A173  1
## 2 A12 48 A32 A43 5951 A61 A73  2 A92 A101  2 A121  22 A143 A152  1 A173  1
## 3 A14 12 A34 A46 2096 A61 A74  2 A93 A101  3 A121  49 A143 A152  1 A172  2
## 4 A11 42 A32 A42 7882 A61 A74  2 A93 A103  4 A122  45 A143 A153  1 A173  2
## 5 A11 24 A33 A40 4870 A61 A73  3 A93 A101  4 A124  53 A143 A153  2 A173  2
## 6 A14 36 A32 A46 9055 A65 A73  2 A93 A101  4 A124  35 A143 A153  1 A172  2
##      V19  V20 V21
## 1 A192 A201  1
## 2 A191 A201  2
## 3 A191 A201  1
## 4 A191 A201  1
## 5 A191 A201  2
## 6 A192 A201  1
```

`str(gcred)`

```
## 'data.frame': 1000 obs. of 21 variables:
## $ V1 : chr "A11" "A12" "A14" "A11" ...
## $ V2 : int 6 48 12 42 24 36 24 36 12 30 ...
## $ V3 : chr "A34" "A32" "A34" "A32" ...
## $ V4 : chr "A43" "A43" "A46" "A42" ...
## $ V5 : int 1169 5951 2096 7882 4870 9055 2835 6948 3059 5234 ...
## $ V6 : chr "A65" "A61" "A61" "A61" ...
## $ V7 : chr "A75" "A73" "A74" "A74" ...
## $ V8 : int 4 2 2 2 3 2 3 2 2 4 ...
## $ V9 : chr "A93" "A92" "A93" "A93" ...
## $ V10: chr "A101" "A101" "A101" "A103" ...
## $ V11: int 4 2 3 4 4 4 4 2 4 2 ...
## $ V12: chr "A121" "A121" "A121" "A122" ...
## $ V13: int 67 22 49 45 53 35 53 35 61 28 ...
## $ V14: chr "A143" "A143" "A143" "A143" ...
## $ V15: chr "A152" "A152" "A152" "A153" ...
## $ V16: int 2 1 1 1 2 1 1 1 1 2 ...
## $ V17: chr "A173" "A173" "A172" "A173" ...
## $ V18: int 1 1 2 2 2 2 1 1 1 1 ...
## $ V19: chr "A192" "A191" "A191" "A191" ...
## $ V20: chr "A201" "A201" "A201" "A201" ...
## $ V21: int 1 2 1 1 2 1 1 1 1 2 ...
```

*#variable 21 contains the responses we want so we need to scale 1 & 2 with 0 & 1*

```
gcred$V21[gcred$V21==1] <- 0
gcred$V21[gcred$V21==2] <- 1
```

*# we want to split the model into training and validation sets so we can compute truepositive + truenegative / (total predictions + actual)*

```
gcredspl <- createDataPartition(gcred$V21, times = 1, p = 0.7, list= F)

head(gcredspl)
```

```
##      Resample1
## [1,]         1
## [2,]         4
## [3,]         5
## [4,]         6
## [5,]         7
## [6,]         8
```

```
#create training and validation sets
```

```
gt <- gcred[gcredspl,]
```

```
gv <- gcred[-gcredspl,]
```

```
table(gt$V21)
```

```
##
```

```
##    0    1
```

```
## 480 220
```

```
table(gv$V21)
```

```
##
```

```
##    0    1
```

```
## 220   80
```

```
#show both sets
```

```
#now we can run the log model with glm
```

```
glogm <- glm(V21~., data = gt, family = binomial(link = "logit"))
```

```
summary(glogm)
```



```
##
## Call:
## glm(formula = V21 ~ ., family = binomial(link = "logit"), data = gt)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9416  -0.7180  -0.3333   0.7217   2.7415
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  7.706e-01  1.291e+00   0.597 0.550461
## V1A12        -6.201e-01  2.715e-01  -2.283 0.022401 *
## V1A13        -1.250e+00  4.730e-01  -2.644 0.008203 **
## V1A14        -1.686e+00  2.806e-01  -6.009 1.87e-09 ***
## V2           3.127e-02  1.124e-02   2.783 0.005385 **
## V3A31         1.550e-01  6.753e-01   0.230 0.818453
## V3A32        -4.830e-01  5.019e-01  -0.962 0.335895
## V3A33        -3.302e-01  5.485e-01  -0.602 0.547138
## V3A34        -1.026e+00  5.074e-01  -2.022 0.043212 *
## V4A41        -2.170e+00  4.890e-01  -4.438 9.09e-06 ***
## V4A410       -3.536e+00  1.190e+00  -2.972 0.002958 **
## V4A42        -8.265e-01  3.183e-01  -2.596 0.009422 **
## V4A43        -8.854e-01  2.984e-01  -2.968 0.003002 **
## V4A44        -3.845e-01  8.537e-01  -0.450 0.652388
## V4A45         1.917e-01  6.398e-01   0.300 0.764486
## V4A46         1.766e-01  4.686e-01   0.377 0.706313
## V4A48        -2.113e+00  1.343e+00  -1.573 0.115676
## V4A49        -9.190e-01  4.045e-01  -2.272 0.023114 *
## V5           1.602e-04  5.432e-05   2.949 0.003188 **
## V6A62        -1.445e-01  3.286e-01  -0.440 0.660123
## V6A63        -8.399e-01  5.249e-01  -1.600 0.109563
## V6A64        -1.476e+00  7.412e-01  -1.991 0.046437 *
## V6A65        -1.063e+00  3.238e-01  -3.283 0.001029 **
## V7A72        -1.147e-02  4.996e-01  -0.023 0.981680
## V7A73        -3.296e-01  4.806e-01  -0.686 0.492899
## V7A74        -5.868e-01  5.205e-01  -1.127 0.259609
## V7A75        -1.608e-01  4.926e-01  -0.326 0.744074
## V8           3.818e-01  1.067e-01   3.577 0.000348 ***
## V9A92        -4.217e-01  4.522e-01  -0.933 0.351025
## V9A93        -1.145e+00  4.433e-01  -2.582 0.009813 **
## V9A94        -4.048e-01  5.389e-01  -0.751 0.452585
## V10A102       8.723e-01  5.433e-01   1.606 0.108359
## V10A103      -1.303e+00  5.426e-01  -2.401 0.016334 *
## V11          -3.881e-02  1.035e-01  -0.375 0.707785
## V12A122       3.513e-01  3.112e-01   1.129 0.259076
## V12A123       3.237e-01  2.835e-01   1.142 0.253499
## V12A124       7.987e-01  4.979e-01   1.604 0.108680
## V13          -1.568e-02  1.169e-02  -1.342 0.179698
## V14A142      -8.240e-01  5.104e-01  -1.615 0.106416
## V14A143      -8.801e-01  2.840e-01  -3.099 0.001943 **
## V15A152      -2.719e-01  2.805e-01  -0.969 0.332378
## V15A153      -5.632e-01  5.853e-01  -0.962 0.335927
```

```
## V16          1.302e-01  2.289e-01   0.569 0.569564
## V17A172      2.240e-01  8.158e-01   0.275 0.783680
## V17A173      3.317e-01  7.781e-01   0.426 0.669939
## V17A174     -4.800e-03  7.942e-01  -0.006 0.995178
## V18          2.604e-01  3.008e-01   0.865 0.386765
## V19A192     -6.539e-02  2.412e-01  -0.271 0.786272
## V20A202     -1.010e+00  7.727e-01  -1.307 0.191262
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 871.48  on 699  degrees of freedom
## Residual deviance: 620.52  on 651  degrees of freedom
## AIC: 718.52
##
## Number of Fisher Scoring iterations: 5
```

*#now we can use our model to predict the number of good/bad and see how good the model is*

```
gcredpred <- predict(glogm, newdata = gv[,-21], type = "response")

table(gv$V21, round(gcredpred))
```

```
##
##      0    1
## 0 193  27
## 1   45  35
```

*#now we can create a confusion matrix so that we can then set a threshold to hopefully get better results i.e. less values that say incorrectly the data is bad credit etc.*

```
gcredpredv2 <- predict(glogm, newdata = gv[,-21], type = "response")
mat <- as.matrix(table(round(gcredpredv2), gv$V21))
names(dimnames(mat)) <- c("Predict", "observe")

mat
```

```
##      observe
## Predict  0    1
##      0 193  45
##      1  27  35
```

```
threshold <- 0.7

mat2 <- as.matrix(table(round(gcredpredv2>threshold), gv$V21))
names(dimnames(mat)) <- c("predict", "observe")

mat2
```

```
##
##      0   1
##  0 212  63
##  1   8  17
```

*#with the threshold we have less incorrect classifications so this is good!*

*#now we can test accuracy*

```
accuracy <- (mat2[1,1]+mat2[2,2])/(mat2[1,1]+mat2[1,2]+mat2[2,1]+mat2[2,2])
accuracy
```

```
## [1] 0.7633333
```

*#so 0.7 is a pretty good threshold as it gives us 72% accuracy*